




# AI Agents That Replace Real Work

## The Future of Automation



|                  |                            |
|------------------|----------------------------|
| <b>Authors:</b>  | Eman Sarhan, Amir Fadila   |
| <b>Course:</b>   | Orchestration of AI Agents |
| <b>Lecturer:</b> | Dr. Yoram Segal            |
| <b>Date:</b>     | June 8, 2026               |

# Contents

|          |                                                              |           |
|----------|--------------------------------------------------------------|-----------|
| <b>1</b> | <b>AI Agents and Real Work</b>                               | <b>1</b>  |
| 1.1      | Introduction: When AI Agents Start Doing Real Work . . . . . | 1         |
| 1.1.1    | What AI Agents Are—and What They Are Not . . . . .           | 1         |
| 1.1.2    | Task-Level Automation, Not Total Job Replacement . . . . .   | 1         |
| 1.1.3    | Human Oversight Remains Central . . . . .                    | 2         |
| 1.1.4    | Preview of This Book . . . . .                               | 2         |
| <b>2</b> | <b>How AI Agents Transform Real Work</b>                     | <b>3</b>  |
| 2.1      | From Human Task to Agentic Workflow . . . . .                | 3         |
| <b>3</b> | <b>Automation Analysis</b>                                   | <b>4</b>  |
| 3.1      | Automation Impact Graph . . . . .                            | 4         |
| 3.2      | Capability Comparison . . . . .                              | 5         |
| 3.3      | Automation Value Model . . . . .                             | 5         |
| 3.3.1    | Using the models in practice . . . . .                       | 6         |
| <b>4</b> | <b>Bilingual Terminology for AI Agent Workflows</b>          | <b>7</b>  |
| 4.1      | Why Bilingual Terminology Matters . . . . .                  | 7         |
| 4.2      | Core English–Hebrew Terminology . . . . .                    | 7         |
| 4.3      | Controlled Mixed-Language Communication . . . . .            | 8         |
| 4.4      | Operational Guidance for Mixed-Language Teams . . . . .      | 8         |
| <b>5</b> | <b>Agents, Work, and Automation</b>                          | <b>9</b>  |
| 5.1      | What Makes AI Agents Different from Chatbots . . . . .       | 9         |
| 5.1.1    | Capability boundaries . . . . .                              | 9         |
| 5.2      | Tool Use and Workflow Automation . . . . .                   | 9         |
| 5.2.1    | Orchestration in practice . . . . .                          | 10        |
| 5.3      | Human-in-the-Loop Supervision . . . . .                      | 10        |
| 5.4      | Risks, Guardrails, and Failure Modes . . . . .               | 10        |
| 5.5      | Business Adoption and the Future of Work . . . . .           | 11        |
| 5.5.1    | Organizational readiness . . . . .                           | 11        |
| 5.6      | Limitations of Full Automation . . . . .                     | 11        |
| 5.7      | Evidence Retention and Audit Trails . . . . .                | 12        |
| 5.7.1    | Minimum evidence package . . . . .                           | 12        |
| 5.8      | Measuring Agent Program Success . . . . .                    | 13        |
| 5.8.1    | Quarterly review questions . . . . .                         | 13        |
| 5.9      | Deployment Checklist for Production Agents . . . . .         | 13        |
|          | <b>References</b>                                            | <b>15</b> |

# Chapter 1

## AI Agents and Real Work

Tool-enabled AI agents are entering operational workflows that once depended entirely on human coordination. This chapter introduces what agents are, how they differ from chatbots, and why task-level automation—with human oversight—is the credible frame for the future of work Exponentia 2026.

### 1.1 Introduction: When AI Agents Start Doing Real Work

Organizations are adopting tool-enabled AI agents to complete work that previously required handoffs between people and software. An agent can retrieve live information, draft outputs, call APIs, and move a task through defined approval steps. That capability is different from a chatbot that answers isolated prompts without owning an outcome Exponentia 2026.

#### 1.1.1 What AI Agents Are—and What They Are Not

AI agents pursue goals across multiple steps inside a workflow. They combine language reasoning with tools: search interfaces, document systems, ticketing platforms, and automation connectors. Workforce research models this as task delegation—workers choose which occupational steps to hand to software and which to keep under direct control *Future of Work with AI Agents: Auditing Automation* 2025. Agents therefore belong in operations design, not only in experimentation labs.

Chatbots optimize conversational fluency. Agents optimize *task completion* within supervised boundaries. The distinction matters when organizations evaluate automation programs: a fluent answer is not the same as a completed research brief, resolved ticket, or approved compliance checklist.

#### 1.1.2 Task-Level Automation, Not Total Job Replacement

Live discussions about “replacing real work” often collapse two different ideas. Task automation removes repeatable steps—data gathering, first-draft writing, routing, and status updates. Role redesign reshapes who owns judgment, approval, and customer relationships. Administrative automation studies show agents absorbing high-volume structured tasks first while strategic decisions remain human-led ReComAI 2026.

This book uses *replacement* in the task sense. Agents can remove manual effort from specific workflows without making credible claims about eliminating entire professions. Practitioner forums emphasize the same boundary: multi-step agents perform well when goals and tools are explicit, but struggle when ambiguity, cross-domain judgment, or incomplete tooling dominates Reddit r/AI\_Agents 2026.

### 1.1.3 Human Oversight Remains Central

Human-in-the-loop review is not a temporary compromise. It is how organizations manage risk as automation scope grows. Operators approve plans, validate evidence, and intercept errors before external actions occur. Auditing frameworks for the future of work ask which tasks employees trust agents to perform and which require persistent human accountability *Future of Work with AI Agents: Auditing Automation 2025*. Sustainable programs encode those tiers as policy—observe, approve, collaborate, or escalate—rather than assuming the model will self-correct under pressure.

### 1.1.4 Preview of This Book

The chapters that follow examine agents and automation from complementary angles:

- **Chapter 2** maps how structured human work becomes an agentic workflow with tools and human review gates.
- **Chapter 3** compares supervision models and quantifies automation value and risk with tables and formulas.
- **Chapter 4** aligns English–Hebrew terminology for mixed-language teams operating agent platforms.
- **Chapter 5** discusses adoption, guardrails, limitations, and realistic measures of program success.

Together, these sections support a practical thesis: AI agents are reshaping how real work gets done by automating bounded tasks inside supervised processes. The opportunity is substantial when organizations match tool-enabled agents to the right work, keep humans on consequential decisions, and measure outcomes by workflow quality rather than hype Exponentia 2026; ReComAI 2026.

## Chapter 2

# How AI Agents Transform Real Work

### 2.1 From Human Task to Agentic Workflow

AI agents do not replace whole jobs at once. They decompose real work into tasks, delegate structured steps to tool-enabled agents, and reserve ambiguous or high-risk decisions for human judgment Exponentia 2026; *Future of Work with AI Agents: Auditing Automation* 2025.

Selective automation is the practical pattern in live research: agents absorb repeatable workflow tasks while organizations keep humans accountable for outcomes that matter ReComAI 2026. Structured tasks can flow through agents with search, APIs, and document tools; complex or underspecified work still needs explicit human decisions Reddit r/AI\_Agents 2026.




Figure 2.1: Agentic automation workflow: structured work can be delegated to AI agents, while high-risk or ambiguous tasks remain under human review.

Agents reshape **how work is done**, not always **who does all of it**. The automation impact graph in Chapter 3 shows where supervision stays high even as structured tasks move to agents *Future of Work with AI Agents: Auditing Automation* 2025.

Workflow commentary argues agents reduce friction between tools and teams rather than replacing every role Exponentia 2026. Administrative automation studies describe task displacement in data handling with role redesign rather than silent team elimination ReComAI 2026. Together, these views support bounded delegation: automate structured work, supervise outcomes, and escalate ambiguous tasks to humans Reddit r/AI\_Agents 2026.

## Chapter 3

# Automation Analysis

### 3.1 Automation Impact Graph

The graph below illustrates workflow maturity, automation impact, and human supervision requirements across task categories *Future of Work with AI Agents: Auditing Automation 2025*.




Figure 3.1: Python-generated graph: automation impact vs human supervision by task type.

Routine administrative work shows the highest automation impact in this model, while complex judgment tasks retain the highest supervision requirement ReComAI 2026. That pattern supports a tiered adoption strategy: automate structured spines first, keep humans on exception handling and approval gates *Future of Work with AI Agents: Auditing Automation 2025*.

The workflow diagram in Chapter 2 explains **how** real work moves through agentic automation. This chart estimates **where** automation pressure and human supervision concentrate in practice.

Table 3.1: Chatbot vs AI Agent vs Agentic Workflow (from live research synthesis).

| Dimension       | Chatbot             | AI Agent                                | Agentic Workflow                        |
|-----------------|---------------------|-----------------------------------------|-----------------------------------------|
| Task scope      | Single-turn replies | Multi-step tasks with tools             | Orchestrated crew <i>Process</i>        |
| Tool use        | Limited plugins     | Live search, APIs, files                | Chained agent <i>Tasks</i> with context |
| Human oversight | Light review        | <i>Human-in-the-loop</i> for edge cases | Supervision gates per stage             |
| Risk profile    | Low automation risk | Medium: bounded autonomy                | Higher: end-to-end automation           |

## 3.2 Capability Comparison

Industry commentary suggests agents reshape workflows rather than replace entire jobs Exponentia 2026; ReComAI 2026. Structured tasks show higher automation potential than complex judgment work Reddit r/AI\_Agents 2026.

## 3.3 Automation Value Model

$$V_{\text{auto}} = T_{\text{saved}} - C_{\text{run}} - R_{\text{risk}}$$

### Where:

- $T_{\text{saved}}$  — time saved by automating repeatable *Tasks*
- $C_{\text{run}}$  — compute, API, and orchestration cost per *Process*
- $R_{\text{risk}}$  — risk cost from errors without *Human-in-the-loop*

A complementary productivity ratio used in workforce audits:

$$\eta_{\text{reduce}} = \frac{N_{\text{automated}}}{N_{\text{total}}} \quad \text{with} \quad 0 \leq \eta_{\text{reduce}} \leq 1$$

When  $\eta_{\text{reduce}}$  rises, required human supervision should be re-evaluated. Academic frameworks for delegating work to agents emphasize auditing automation boundaries *Future of Work with AI Agents: Auditing Automation* 2025.

### 3.3.1 Using the models in practice

Estimate  $T_{\text{saved}}$  from timed workflows before automation, not from vendor demos. Record baseline minutes per research brief, draft cycle, and review pass; after deployment, compare distributions rather than single runs ReComAI 2026. Estimate  $C_{\text{run}}$  from token logs, search costs, and human review time Exponentia 2026. Estimate  $R_{\text{risk}}$  from incident classes—bad citations, leaked secrets, compliance exceptions—and cap  $\eta_{\text{reduce}}$  until measurement improves Reddit r/Al\_Agents 2026. A program is healthy when  $V_{\text{auto}}$  stays positive *after* supervision and audit work is included.



## Chapter 4

# Bilingual Terminology for AI Agent Workflows

AI-agent workflows are often designed in English, while local teams may document decisions, approvals, and operational policy in Hebrew. This chapter defines shared terminology for automation, supervision, and real work across both languages Exponentia 2026.

### 4.1 Why Bilingual Terminology Matters

English dominates vendor documentation, model APIs, and agent-framework vocabulary. Yet organizations that automate real work must explain credentials, review gates, and escalation rules to Hebrew-speaking operators and managers. Without aligned terms, teams misread what an *Agent*, *Task*, or *Process* actually authorizes in production *Future of Work with AI Agents: Auditing Automation* 2025.

Live workforce research asks which occupational tasks workers are willing to delegate to agents versus keep under direct human control *Future of Work with AI Agents: Auditing Automation* 2025. Bilingual terminology is therefore not cosmetic—it is how mixed-language teams agree on boundaries before automation scales ReComAI 2026.

### 4.2 Core English–Hebrew Terminology

Each definition below lists the English term, Hebrew equivalent, and Meaning in workflow. Definition rows avoid tabular BiDi conflicts while keeping bilingual pairs readable.

- **Agent** — סוכן — Executes a bounded task using tools.
- **Task** — משימה — A unit of work delegated to an agent.
- **Crew** — צוות — A coordinated group of agents or roles.
- **Process** — תהליך — Ordered workflow from input to output.
- **API Key** — מפתח גישה — Credential used to access external services.
- **Human-in-the-loop** — בקרה אנושית — Human review before high-risk output.

## 4.3 Controlled Mixed-Language Communication

**Mixed-language example.** In an Israeli engineering team, a design review might use English terms such as Agent, Task, API Key, and Human-in-the-loop, while the operational policy is explained in Hebrew.

The English labels above are standard in agent platforms. The operational policy itself can be stated in Hebrew without embedding those Latin tokens inside Hebrew sentences.

בדיון פנימי, הצוות מפרט את מדיניות הפעולה בעברית.

Practitioner forums note that complex delegation still fails when boundaries are vague [Reddit r/AI\\_Agents 2026](#). Workflow commentary stresses agents as connectors between tools and teams—not silent replacements for accountability [Exponentia 2026](#).

### סיכום בעברית

סוכני בינה מלאכותית יכולים לבצע משימות מוגדרות כחלק מתהליך עבודה מתוזמר. במערכת כזו, המשימה מפורקת לשלבים קטנים וברורים. הסוכן משתמש בכלים כדי לאסוף מידע, ליצור טיוטה או לבצע פעולה. כאשר המשימה רגישה, מסוכנת או לא ברורה, נדרשת בקרה אנושית. המטרה אינה להחליף בני אדם לגמרי, אלא להעביר משימות מובנות לאוטומציה תוך שמירה על אחריות, פיקוח ואיכות. לכן חשוב להגדיר מראש מתי הסוכן פועל לבד ומתי האדם חייב לאשר את התוצאה.

## 4.4 Operational Guidance for Mixed-Language Teams

When Hebrew-English teams deploy agents into real work, three practices keep automation safe:

- Publish the terminology definitions in design reviews so every role shares the same workflow vocabulary.
- Store credentials and API access rules in a single authoritative runbook language, with Hebrew summaries where operators require them.
- Mark customer-facing and compliance-sensitive steps as human-in-the-loop regardless of language [Reddit r/AI\\_Agents 2026](#); [ReComAI 2026](#).

Automation can accelerate structured tasks in any language context, but approval language must be explicit before agents act on consequential work [Future of Work with AI Agents: Auditing Automation 2025](#).

## Chapter 5

# Agents, Work, and Automation

This chapter synthesizes live internet research on how AI agents change real work, workflow automation, and organizational adoption across industries entering the agent era.

### 5.1 What Makes AI Agents Different from Chatbots

Chatbots answer prompts in a single conversational turn. AI agents pursue goals across multiple steps: they search, call tools, write drafts, and pass context to downstream tasks. Live commentary frames this shift as workflow orchestration rather than simple dialogue Exponentia 2026.

A chatbot optimizes for fluent replies. An agent optimizes for *task completion* inside a bounded process with tools, memory, and supervision gates.

Agents connect signals across systems—email, tickets, documents, APIs—and reduce hand-offs that slow teams. Production teams therefore separate research, drafting, and quality review instead of relying on one monolithic prompt Exponentia 2026.

Administrative automation reports show agents absorbing routine data entry and document extraction first—high-volume, structured work ReComAI 2026. Chatbots rarely own end-to-end outcomes in those settings because they lack persistent tool access and review checkpoints.

#### 5.1.1 Capability boundaries

Community practitioners note that agents excel within well-defined boundaries but struggle when judgment, ambiguity, or cross-domain reasoning dominate Reddit r/AI\_Agents 2026. The gap between “automate a workflow step” and “replace a profession” remains wide in live discussions.

### 5.2 Tool Use and Workflow Automation

Modern agent stacks treat tools as first-class interfaces: search services, document repositories, ticketing systems, and enterprise API connectors. Each tool invocation is a *Task* with an auditable input and output, which supports traceability in regulated environments.

Typical pattern: **search** → **synthesize** → **draft** → **review** → **publish**. Each stage can fail independently and be retried without corrupting upstream evidence.

Workflow automation value rises when agents observe events, prioritize work, and route tasks between humans and software Exponentia 2026. Tool use is therefore not a feature list—it is the mechanism that turns language models into operators.

Academic auditing frameworks ask which occupational tasks workers would delegate to agents versus keep under human control *Future of Work with AI Agents: Auditing Automation* 2025. Tool-enabled agents make those trade-offs explicit: every automated step should map to a named capability and a supervision policy.

### 5.2.1 Orchestration in practice

Sequential orchestration runs specialized agents with shared context. That design mirrors how organizations already separate research, drafting, and quality review—but at higher speed and with live source retrieval. Chapter 2 illustrates how structured tasks flow through agent automation; Chapter 3 quantifies where automation impact is highest across supervision models.

## 5.3 Human-in-the-Loop Supervision

*Human-in-the-loop* review is not optional for high-impact automation. Operators approve agent plans, validate citations, and intercept errors before publication or customer-facing actions. Research on occupational delegation stresses auditing which tasks workers trust agents to perform *Future of Work with AI Agents: Auditing Automation* 2025.

Increase automation scope only when monitoring cost stays lower than the value of delegated work. If review queues grow faster than throughput, the agent design is likely too autonomous for the risk profile.

Supervision tiers can be documented as:

- **Observe:** human watches logs and metrics only.
- **Approve:** human confirms before external side effects.
- **Collaborate:** human and agent iterate on the same task.
- **Escalate:** agent hands off when confidence or policy thresholds fail.

Reddit practitioners highlight that multi-step agents need clear stop conditions when tasks exceed training or tool boundaries Reddit r/AI\_Agents 2026. Production systems should encode those stops as policy, not hope the model improvises safely.

## 5.4 Risks, Guardrails, and Failure Modes

Agent failures differ from chatbot failures. A wrong answer becomes a wrong action: bad citations, leaked secrets, duplicated customer emails, or corrupted documents. Guardrails therefore span data access, tool permissions, and output validation—not just prompt engineering.

- Hallucinated sources that look plausible in draft text.

- Over-automation of judgment-heavy decisions.
- Credential exposure through logs or prompts.
- Runaway tool loops without budget or time limits.

Mitigations align with the automation value model in Chapter 3: track runtime cost  $C_{run}$ , estimate risk cost  $R_{risk}$ , and require human gates when expected loss exceeds savings. Administrative automation case studies show large gains in routine roles but also role redesign rather than silent elimination ReComAI 2026.

## 5.5 Business Adoption and the Future of Work

Enterprises adopt agents where friction is measurable: ticket triage, research briefs, compliance checklists, and document pipelines. The near-term story is less “fewer jobs” and more “fewer manual handoffs”—agents sit between teams and systems Exponentia 2026.

Measure adoption by **workflow latency** and **error rework**, not only headcount. Agents that shorten review cycles while improving citation quality justify continued investment even when full replacement is unrealistic.

Workforce auditing research proposes structured evaluation of which tasks employees want automated versus augmented *Future of Work with AI Agents: Auditing Automation* 2025. That employee-centered framing reduces backlash and produces guardrails grounded in real job constraints.

Industry reporting on administrative roles shows concrete task displacement patterns—data entry, scheduling support, extraction—while strategic decisions remain human-led ReComAI 2026. Future-of-work narratives should separate *task automation* from *role erasure*.

### 5.5.1 Organizational readiness

Readiness checklist for agent deployment:

- Live source policy: which APIs and search tools are approved?
- Secret management: keys outside prompts and version control.
- Review ownership: who signs drafts before external release?
- Evidence retention: store search results and agent logs for audit.
- Rollback plan: how to revert automated changes quickly.

## 5.6 Limitations of Full Automation

Full automation—agents operating without meaningful human oversight—is rarely credible for complex, high-stakes work. Practitioners report agents handling multi-step routines but failing when goals are underspecified or tools are incomplete Reddit r/AI\_Agents 2026.

Agents are force multipliers for structured work, not universal substitutes for expertise. The central theme—agents that replace *real work*—must be read as selective automation of tasks, not total workforce replacement.

Limitations cluster into four areas:

1. **Context limits:** agents lose thread across long, branching projects.
2. **Judgment limits:** ethical, legal, and brand decisions need human owners.
3. **Tool limits:** agents cannot act where APIs or data are unavailable.
4. **Verification limits:** citations and numbers require independent checks *Future of Work with AI Agents: Auditing Automation* 2025.

Sustainable programs treat agents as teammates in a supervised *Process*: automate the repetitive spine of work, keep humans on the margins where variance and accountability concentrate. That posture matches both academic auditing models and practitioner caution in live forums Reddit r/AI\_Agents 2026; Exponentia 2026.

## 5.7 Evidence Retention and Audit Trails

Automated research and drafting are only defensible when organizations retain evidence. Each agent stage should persist inputs and outputs: source records, intermediate drafts, approval notes, and published deliverables. That mirrors how compliance teams reconstruct decisions after incidents *Future of Work with AI Agents: Auditing Automation* 2025.

Store **who** triggered a run, **which tools** were called, **what sources** were retrieved, and **who approved** external publication. Agents without audit trails are difficult to defend in enterprise settings.

Audit trails also reduce citation hallucination risk. When reviewers can compare draft claims to stored source snippets, false references surface before publication. Administrative automation programs report the same pattern: automation succeeds when document provenance is explicit ReComAI 2026.

### 5.7.1 Minimum evidence package

- Timestamped source records with URLs and titles.
- Intermediate drafts with role labels (research, draft, approval).
- Review outcomes covering accuracy, clarity, and policy alignment.
- Published outputs with version identifiers.
- Runtime and cost summaries for budget governance.

## 5.8 Measuring Agent Program Success

Organizations should score agent programs with operational metrics, not hype metrics. Vanity metrics—“number of prompts” or “model upgrades”—do not prove work was replaced safely.

Table 5.1: Sample metrics for agent automation programs.

| Metric           | What it measures                       | Why it matters          |
|------------------|----------------------------------------|-------------------------|
| Cycle time       | Hours from research to approved draft  | Workflow speed          |
| Rework rate      | Share of drafts sent back for revision | Quality and supervision |
| Source coverage  | Citations tied to live URLs            | Research integrity      |
| Escalation rate  | Tasks routed to humans                 | Boundary respect        |
| Cost per outcome | API spend per published artifact       | Financial guardrail     |

Academic workforce auditing suggests measuring *which tasks workers want delegated* rather than assuming universal enthusiasm for automation *Future of Work with AI Agents: Auditing Automation 2025*. Pair that qualitative signal with the quantitative table above.

Do not reward agents for volume alone. Publishing more drafts with weak citations increases rework and reputational risk faster than it saves labor Reddit r/AI\_Agents 2026.

### 5.8.1 Quarterly review questions

1. Which workflows gained measurable latency reduction?
2. Where did human review time increase instead of decrease?
3. Did live source diversity improve or narrow?
4. Which tool permissions should be tightened?
5. Are high-risk tasks still behind human approval gates?

Answering these questions keeps agent programs aligned with the book’s central theme: **AI agents that replace real work**—meaning specific, evidenced tasks—rather than unbounded promises of full workforce replacement Exponentia 2026; ReComAI 2026.

## 5.9 Deployment Checklist for Production Agents

Before running agents against live APIs in production, teams should verify environment, governance, and rollback mechanics. The checklist below summarizes practices supported by live research on workflow agents Exponentia 2026; *Future of Work with AI Agents: Auditing Automation 2025*.

1. **Govern credentials:** store API keys and tokens outside prompts and version-controlled content.
2. **Validate tools:** confirm connectors return authoritative data before drafting or executing actions.

3. **Scope rollouts:** start with one workflow; expand only after latency and rework metrics improve ReComAI 2026.
4. **Validate outputs:** ensure drafts and citations match retrieved evidence.
5. **Maintain references:** linked citations must map to a curated bibliography.
6. **Archive artifacts:** retain source records, drafts, approvals, and audit logs.
7. **Review high-risk text:** require human sign-off when outputs mention job replacement or unverified claims Reddit r/AI\_Agents 2026.

Treat each agent pipeline execution as a **publishable workflow**, not a demo. The difference is evidence: stored artifacts, linked citations, and reproducible outputs.

Administrative automation case studies emphasize phased rollout: automate one workflow, measure rework and latency, then expand tool permissions gradually ReComAI 2026. Skipping phases produces brittle programs that break when models, APIs, or policies change.

Finally, communicate limits openly. Agents that augment workflows can earn trust; agents marketed as flawless replacements trigger resistance and hidden manual work Exponentia 2026; Reddit r/AI\_Agents 2026. Honest limitation language is part of professional design—not a weakness in the automation story.



# References

- Exponentia (2026). *AI Agents Are Not the Future of Work, They're the Future of Workflows*. Retrieved during live web research for this project. URL: <https://medium.com/@exponentia.ai/ai-agents-are-not-the-future-of-work-theyre-the-future-of-workflows-41af96fc2554>.
- Future of Work with AI Agents: Auditing Automation* (2025). Retrieved during live web research for this project. arXiv: 2506.06576 [cs.AI]. URL: <https://arxiv.org/abs/2506.06576>.
- ReComAI (2026). *How AI Agents Are Replacing Job Roles in Administrative Work*. Retrieved during live web research for this project. URL: <https://www.recomai.one/ai-insights-hub/ai-agents-replacing-job-roles-2026>.
- Reddit r/AI\_Agents (2026). *Can AI Agents really replace humans for complex tasks?* Retrieved during live web research for this project. URL: [https://www.reddit.com/r/AI\\_Agents/comments/1s2d80f/can\\_ai\\_agents\\_really\\_replace\\_humans\\_for\\_complex/](https://www.reddit.com/r/AI_Agents/comments/1s2d80f/can_ai_agents_really_replace_humans_for_complex/).